

Adversarial Instance Re-weighting for Unsupervised Domain Adaptation

Itzel Belderbos

29 Oct, 2021

DKE supervisors:
Drs. Mirela Popa
Dr. Christof Seiler

External supervisor:
Dr. Fabian Jansen

Contents

1. Introduction

1. Problem statement
2. Research goal & questions

2. Part 1: proof-of-concept

1. Methodology
2. Results

3. Part 2: state-of-the-art comparison

1. Methodology
2. Results

4. Discussion

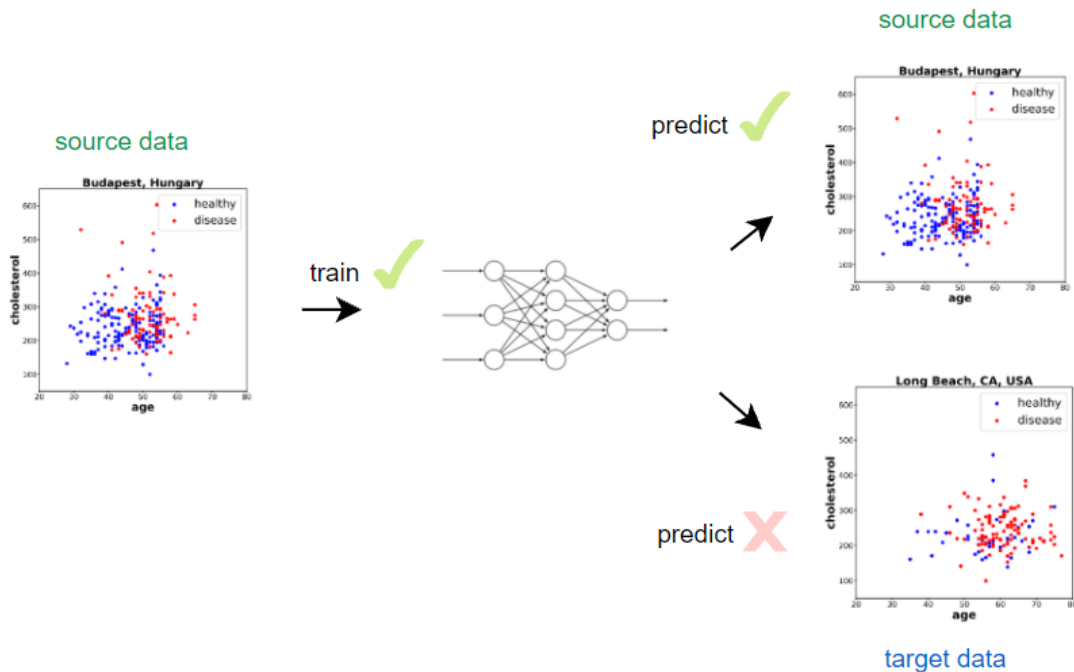
5. Conclusion

Problem statement

- **Wanted:**
 - Train model on *source* dataset, test on *target* dataset
- **Problem:**
 - Target dataset should have same distribution!

→ Labels target dataset not available

→ Data annotation is expensive



Examples

- Computer vision
 - From one weather condition to another
- Sentiment analysis
 - From book reviews to DVD reviews
- Speech recognition
 - From one speaker to another speaker

→ Focus in this work on **images**



Domain adaptation

- Domain: a distribution within the feature space
- This work focuses to two types of DA:

Covariate shift

$$\begin{aligned}\mathcal{X}^s &= \mathcal{X}^t \\ P(Y|X)^s &= P(Y|X)^t \\ P(X)^s &\neq P(X)^t\end{aligned}$$

Prior probability shift

$$\begin{aligned}\mathcal{X}^s &= \mathcal{X}^t \\ P(X|Y)^s &= P(X|Y)^t \\ P(Y)^s &\neq P(Y)^t\end{aligned}$$

Research goal

- Introduce *weight network* to transform distribution sample-wise

→ **Goal**

Investigate to which extent this weight network is able to benefit the domain adaptation, particularly for covariate shift

Research questions

1. Can we use non-generative adversarial networks to improve covariate shift adaptation?
 - a. How can we train a weight function, which converts source into target distribution?
 - b. What is the effect of the critic?
1. How can we measure the performance of the reconstructed target distribution and the adversarial network?
1. Are there other ways than a weight function to express source data into target data?

Methodology for proof-of-concept

The Wasserstein GAN vs our network

- Wasserstein GAN
 - Adversarial game between critic & generator

$$\min_g \max_c \mathbb{E}_{x \sim p_{data}(x)} [c(x)] - \mathbb{E}_{z \sim p_z(z)} [c(g(z))]$$

- Our weighted network
 - Adversarial game between critic & weight

$$\min_w \max_c \mathbb{E}_{x \sim s(x)} [c(x)w(x)] - \mathbb{E}_{x \sim t(x)} [c(x)]$$

Experiments

- **Experiment 1:** One-dimensional Gaussians
- **Experiment 2:** MNIST subset & class imbalance
- **Experiment 3:** MNIST subset & covariate shift

Evaluation metric

→ **Binary domain classifier** to evaluate *domain invariance*

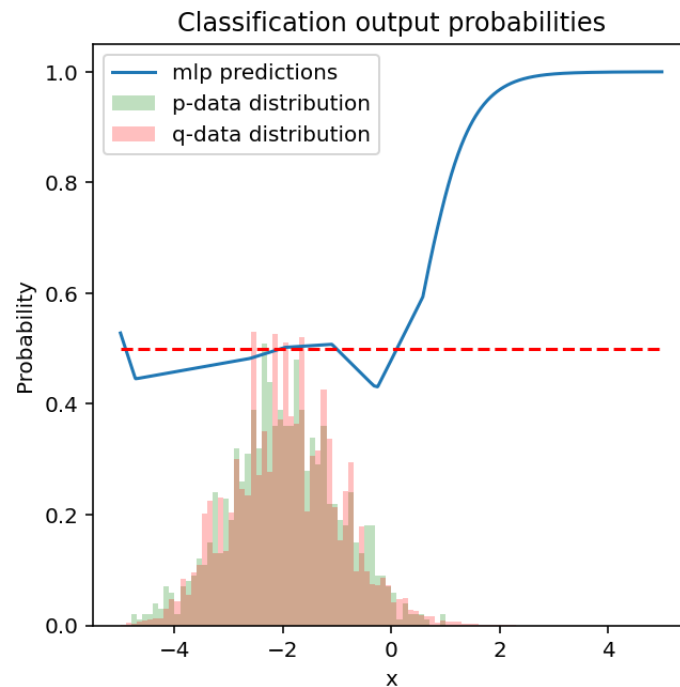
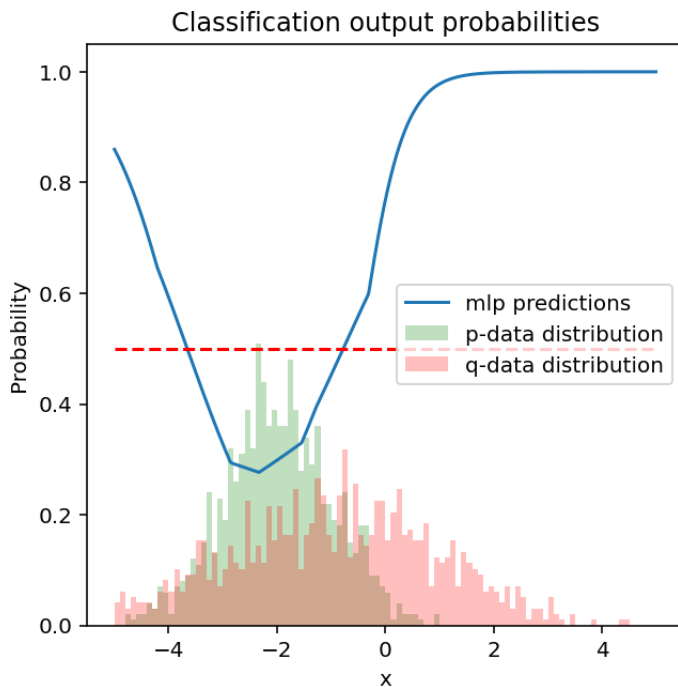
- Combine source and target dataset into 1 dataset
- Label target as 0, source as 1
- Check whether classifier can see difference
- Ideal situation: accuracy of 50%

Results for proof-of-concept

Experiment 1 - Results

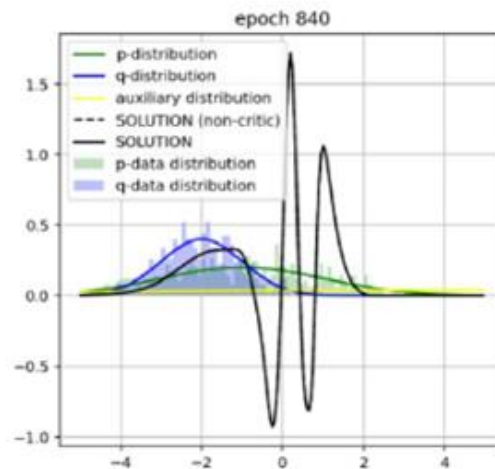
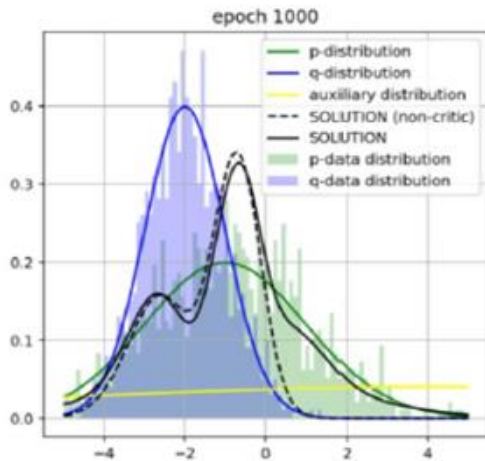
- Weighted distribution: binary domain classifier cannot see difference

green = target
red = source



Experiment 1 - Limitations

- Target dataset cannot contain values which are not present in source dataset
 - Cannot re-weight values which do not exist



Experiment 2 - Prior probability shift

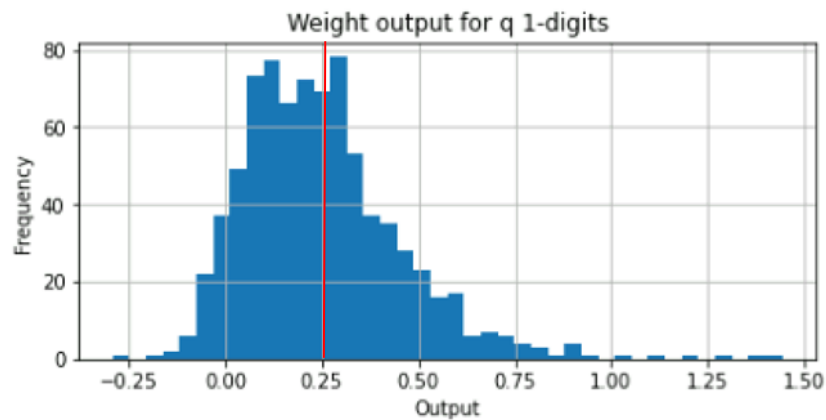
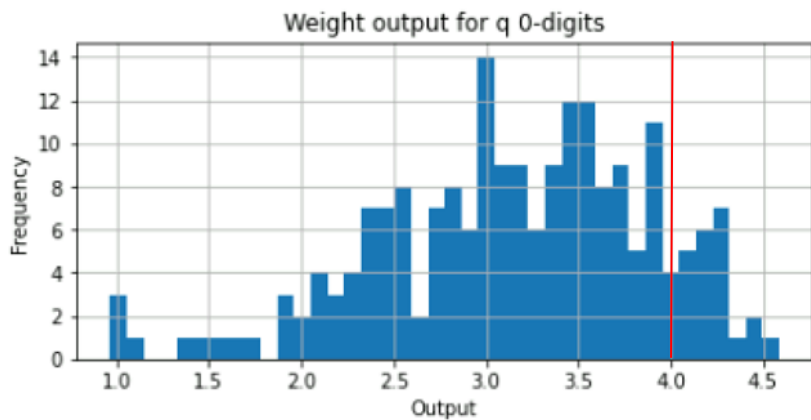
- 1000 samples; classes 0 and 1
- Imbalance between 2 classes

k	$P(Y = k)^s$	$P(Y = k)^t$	$\mathbb{E}[w(x)]$
0	0.20	0.80	4
1	0.80	0.20	$\frac{1}{4}$

Experiment 2 - Results


- Histogram of weight outputs

k	$P(Y = k)^s$	$P(Y = k)^t$	$\mathbb{E}[w(x)]$
0	0.20	0.80	4
1	0.80	0.20	$\frac{1}{4}$



Experiment 2 - Results

- Binary domain classifier
 - **x-axis:** epochs
 - **y-axis:** accuracy



Experiment 3 - Covariate shift

- Source/target datasets consist of 10,000 samples with class 0-4
- Use KMeans to obtain 2 **subclasses** per class
- Create imbalance between 2 subclasses

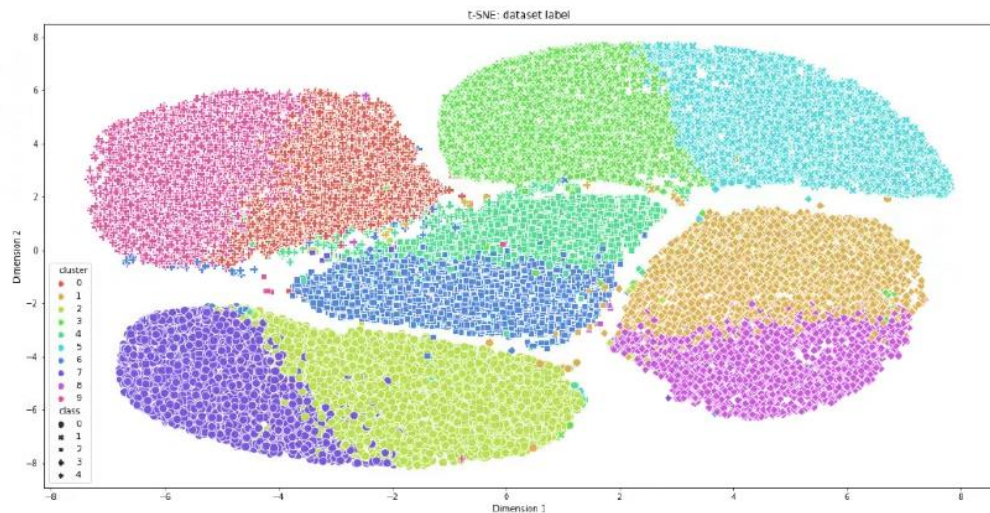




Figure 3.17: t-SNE of subset with 5 classes

Experiment 3 - Results

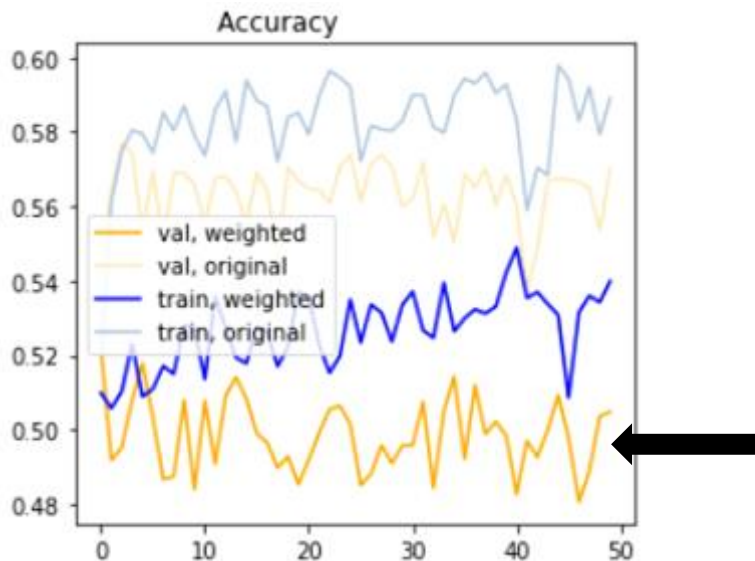


class	cluster	expected weight	actual weight
0	2	1	0.97
0	7	1	0.98
1	3	1	0.98
1	5	1	0.99
2	4	1	0.90
2	6	1	0.99
3	0	0.33	0.55
3	9	3	1.92
4	1	0.33	0.56
4	8	3	2.03


Table 3.3: Expected weight vs actual weight per cluster

Experiment 3 - Results

- Binary domain classifier


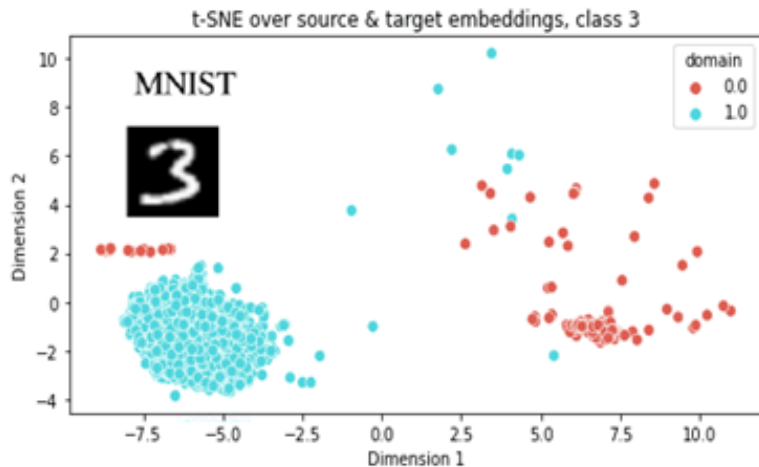


Experiment 3 - Results



Methodology for state-of-the-art comparison

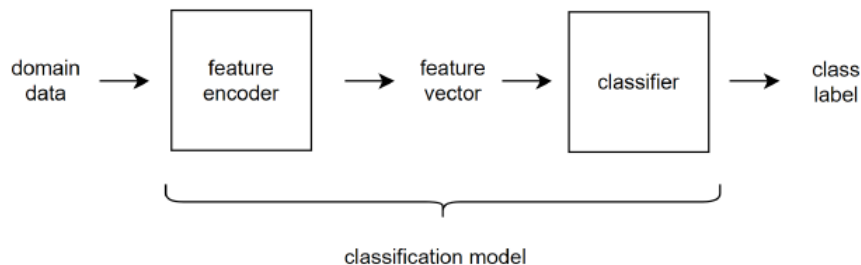
Large domain shift: different features



What to do when **target features** are not present in source data?

Combination with feature-based domain adaptation

- Learn a domain-invariant feature embedding with a **feature encoder**



$$P(X^s) \neq P(X^t)$$



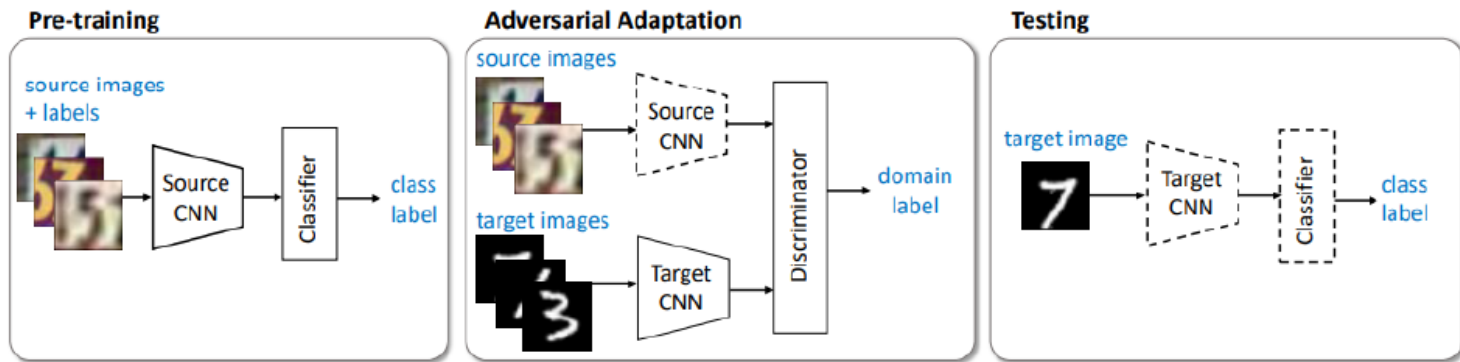
$$P(\tau(X^s)) = P(\tau(X^t))$$

Baseline models

- Adversarial Discriminative Domain Adaptation (ADDA)
- Wasserstein Guided Representation Learning framework (WDGRL)
 - Adversarial game between **discriminator/critic** & **feature encoder**

ADDA

- Works like a GAN, instead of *fake image* it learns a *fake feature representation*
- Binary cross-entropy loss
- 2 separate encoders for **source** and **target** (asymmetric)



ADDA

- Works like a GAN, instead of *fake image* it learns a *fake feature representation*
- Binary cross-entropy loss
- 2 separate encoders for **source** and **target** (asymmetric)

discriminator	$\min_d L^d = -\mathbb{E}[\log d(\tau^s(x^s))] - \mathbb{E}[\log(1 - d(\tau^t(x^t)))]$
target encoder	$\min_{\tau^t} L^{\tau^t} = -\mathbb{E}[\log d(\tau^t(x^t))]$

WDGRL

- Uses Wasserstein distance instead of BCE loss
- 1 encoder for both **source** and **target** (*symmetric*)
- Adds **classification loss** in feature loss
 - Feature loss = Wasserstein loss + source classification loss

↑ ↑
domain invariance class segregation

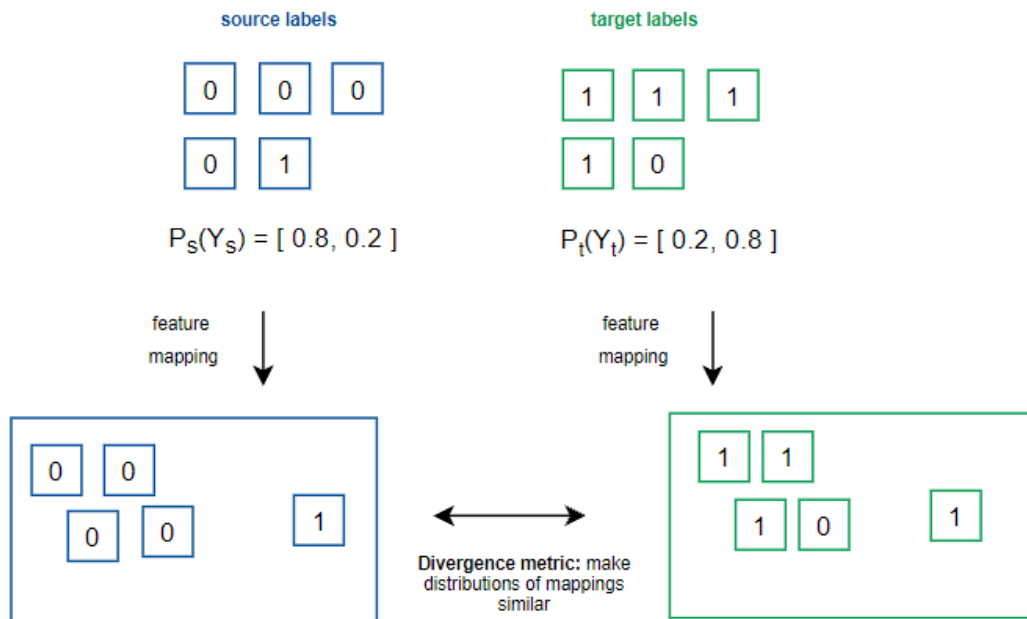
WDGRL

- Uses Wasserstein distance instead of BCE loss
- 1 encoder for both **source** and **target** (symmetric)
- Adds **classification loss** in feature loss
 - Feature loss = Wasserstein loss + source classification loss

critic	$\min_d L^d = \underbrace{\frac{1}{T} \sum_{x \in X^t} d(\tau^{s,t}(x^t)) - \frac{1}{S} \sum_{x \in X^s} d(\tau^{s,t}(x^s))}_{\text{Wasserstein loss}} + \underbrace{\nu \frac{1}{E} \sum_{x \in X^s} (\ \nabla_e d\ _2 - 1)^2}_{\text{gradient penalty}}$
feature encoder	$\min_{\tau^{s,t}, c} L^{ce, wdgrl} L^\tau = \lambda \underbrace{\left[\frac{1}{S} \sum_{x \in X^s} d(\tau^{s,t}(x^s)) - \frac{1}{T} \sum_{x \in X^t} d(\tau^{s,t}(x^t)) \right]}_{\text{Wasserstein loss}} + \underbrace{L^{ce, wdgrl}}_{\text{classification loss}}$ <p style="text-align: center; margin-top: -10px;">balancing parameter</p>

Introduction of weight network

- Regular domain adaptation methods assume class balance
- Mapping between target data and representations might go wrong



Proposed method

- Addition of **weight**
- Combination of ADDA and WDGRL
 - **WDGRL**: Wasserstein distance & addition of **classification loss** in feature loss
 - **ADDA**: 2 separate encoders for **source** and **target**

Proposed method


- 2 separate encoders for **source** and **target** (ADDA)
- Addition of classification loss in feature loss (WDGRL)
- Addition of **weight**

critic	$\min_d L^d = \underbrace{\frac{1}{T} \sum_{x \in X^t} d(\tau^t(x^t)) - \frac{1}{S} \sum_{x \in X^s} d(\tau^s(x^s)) w(\tau^s(x^s))}_{\text{Wasserstein loss}} + \underbrace{\nu \frac{1}{E} \sum_{x \in X^e} (\ \nabla_e d(e)\ _2 - 1)^2}_{\text{gradient penalty}}$
feature encoder	$\min_{\tau^t} L^{\tau} = \lambda^{\tau^t} \left[\underbrace{-\frac{1}{T} \sum_{x \in X^t} d(\tau^t(x^t))}_{\text{Wasserstein loss}} \right] + \underbrace{L^{ce,ours}}_{\text{classification loss}}$
weight	$\min_w L^w = \lambda^w \left[\underbrace{\frac{1}{S} \sum_{x \in X^s} d(\tau^s(x^s)) w(\tau^s(x^s))}_{\text{Wasserstein loss}} \right] + \underbrace{L^{ce,ours}}_{\text{classification loss}}$

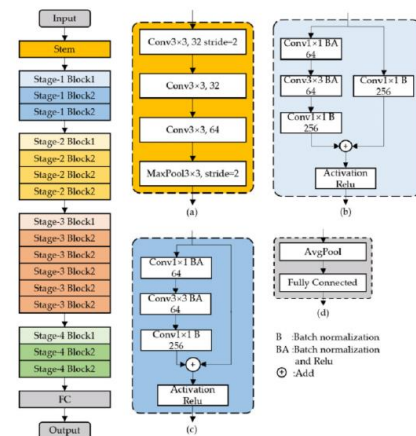
Datasets & backbone architectures



modified LeNet



ResNet-50



Experiments

- **Experiment 4:** unweighted feature-based adaptation
- **Experiment 5:** weighted feature-based adaptation
- **Experiment 6:** weighted feature-based *imbalanced* adaptation

Results for state-of-the-art comparison

Experiment 4: unweighted feature-based model

- MNIST-USPS

Method		$M \rightarrow U$	$U \rightarrow M$
Source only		0.761	0.588
Previous work	DANN [7]	0.771	0.750
	DDC [16]	0.791	0.698
	ADDA [1]	0.896	0.909
	WDGRL [6]	-	-
	CoGAN [8]	0.912	0.899
	DRAnet [11]	0.978	0.991
This work	WDGRL ^{ours}	0.920	0.898
	ADDA ^{ours}	0.901	0.923
	ours	0.958	0.936

Experiment 4: unweighted feature-based model

- MNIST-USPS

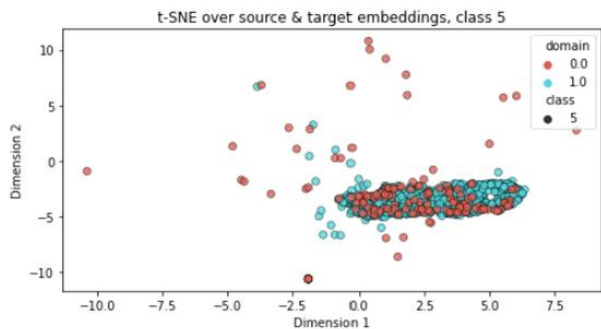


Figure 5.3: Example feature representation for ADDA




Figure 5.4: Example feature representation for WDGRL






Figure 5.5: Example feature representation for our method

Experiment 4: unweighted feature-based model

- Office-31

	Method	$A \rightarrow W$	$D \rightarrow W$	$W \rightarrow A$
	Source only	0.626	0.961	0.610
Previous work	DANN [7]	0.730	0.964	0.675
	ADDA [1]	0.751	0.970	-
	DDC [16]	0.618	0.950	-
	WDGRL [6]	0.895	0.979	0.937
	DADA [18]	0.924	0.993	0.743
This work	ADDA ^{ours}	0.765	0.972	0.697
	ours	0.803	0.980	0.709

Experiment 4: Office-31



Experiment 5: integration with weight

- Large design choice space

critic	$\min_d L^d = \frac{1}{T} \sum_{x \in X^t} d(\tau^t(x^t)) - \frac{1}{S} \sum_{x \in X^s} d(\tau(x^s))w(\tau^s(x^s)) + \nu \frac{1}{E} \sum_{x \in X^e} (\ \nabla_e d(e)\ _2 - 1)^2$
feature encoder	$\min_{\tau^t} L^\tau = \boxed{\lambda^{\tau^t}} \left[-\frac{1}{T} \sum_{x \in X^t} d(\tau^t(x^t)) \right] + \boxed{L^{cc,ours}}$
weight	$\min_w L^w = \boxed{\lambda^w} \left[\frac{1}{S} \sum_{x \in X^s} d(\tau^s(x^s))w(\tau^s(x^s)) \right] + \boxed{L^{cc,ours}}$


Experiment 5: integration with weight

Method	λ_τ	λ_w	$L^{ce,ours}$ for τ	$L^{ce,ours}$ for w	w in $L^{ce,ours}$	w_{pass}	activation	acc_{target}
Source-only								0.761
Unweighted _{ours}								0.958
trial 1	✓	✓	✓	×	✓	×	softmax	0.834
trial 2	✓	×	✓	×	✓	×	softmax	0.818
trial 3	×	×	✓	✓	✓	×	softmax	0.792
trial 4	✓	×	✓	✓	✓	×	relu	0.801
trial 5	✓	×	✓	×	✓	×	relu	0.799
trial 6	✓	✓	✓	✓	✓	×	softmax	0.842
trial 7	✓	✓	✓	✓	✓	✓	softmax	0.862


Table 5.3: Results for weighted domain adaptation on MNIST-USPS transfer

Experiment 6: imbalanced domain adaptation


- **x-axis:** class
- **y-axis:** probability



KL = 0.04



KL = 0.40



KL = 0.74

Experiment 6: imbalanced domain adaptation

Method	$M \rightarrow U$		
$KL(S, T)$	0.04	0.40	0.74
Unweighted	0.872	0.609	0.640
Weighted	0.865	0.714	0.643

Discussion & Conclusion

- Research focus
 - Investigate whether an adversarial instance re-weighting framework, inspired by Wasserstein GAN, is able to enhance unsupervised domain adaptation problem
- Proof-of-concept: small domain shift
 - **Experiment 1:** one-dimensional Gaussians
 - Source and target domain should have sufficient overlap
 - **Experiment 2:** MNIST, class imbalance
 - Algorithm is able to re-weight based on class
 - **Experiment 3:** MNIST, covariate shift
 - Algorithm is able to re-weight based on features

Discussion & Conclusion

- Comparison with state-of-the-art: large domain shift
 - **Experiment 4:** unweighted feature-based adaptation
 - Our method improved ADDA and WDGRL for MNIST-USPS and Office-31
 - Did not improve state-of-the-art
 - **Experiment 5:** weighted feature-based adaptation
 - Large choice space leads to difficult optimization
 - Improves source-only; does not improve unweighted methods
 - **Experiment 6:** weighted feature-based & imbalanced adaptation
 - For a KL-divergence of 0.40 between MNIST and USPS, weighted method outperforms unweighted method